

# Automating Data Exploration with R

## Reading Data

## readLines

A first investigative exploration of a data set is the `readLines` function (<https://stat.ethz.ch/R-manual/R-devel/library/base/html/readLines.html>). It allows you to cull a small amount of lines from the top of a file, no matter how big the file is. This is a silly example but imagine if your file is over 10 gigabytes in size, why waste time and memory when you aren't sure what it contains?

Let's use the `readLines` function to open part of a text file off **CRAN**:

```
readLines('https://cran.r-project.org/src/base/README', n=20)
```

```
## [1] ""
## [2] "\\t\\t\\tTHE BASIC R README"
## [3] ""
## [4] ""
## [5] "\\t    (See \\\"doc/FAQ\\\" and \\\"doc/RESOURCES\\\" for more detailed informatio
n"
## [6] "\\t\\t\\t\\t    - these files are only in the tarballs)"
## [7] "\\t    (See \\\"INSTALL\\\"                for help on installation)"
## [8] ""
## [9] "1. INTRODUCTION"
## [10] ""
## [11] "This directory contains the source code tree for R, which is a"
## [12] "language which is not entirely unlike (versions 3 and 4 of) the S"
## [13] "language developed at AT&T Bell Laboratories by Rick Becker, John"
## [14] "Chambers and Allan Wilks."
## [15] ""
## [16] "R is free software distributed under a GNU-style copyleft."
## [17] ""
## [18] "The core of R is an interpreted computer language with a syntax"
## [19] "superficially similar to C, but which is actually a \\\"functional"
## [20] "programming language\\\" with capabilities similar to Scheme.  The"
```

This is a great way of discovering what data types are contained in a very large data set. You will see in the next few functions we use, knowing the data type in advance can speed things up tremendously.

## read.table, read.csv, read.csv2

`read.csv` is the most common reader in R, followed closely by `read.table`. In fact, `read.csv` and `read.csv2` are both wrappers over `read.table` (<https://stat.ethz.ch/R-manual/R-devel/library/utils/html/read.table.html>).

```
read.table(file, header = FALSE, sep = "", quote = "\"",
           dec = ".", numerals = c("allow.loss", "warn.loss", "no.loss"),
           row.names, col.names, as.is = !stringsAsFactors,
           na.strings = "NA", colClasses = NA, nrows = -1,
           skip = 0, check.names = TRUE, fill = !blank.lines.skip,
           strip.white = FALSE, blank.lines.skip = TRUE,
           comment.char = "#",
           allowEscapes = FALSE, flush = FALSE,
           stringsAsFactors = default.stringsAsFactors(),
           fileEncoding = "", encoding = "unknown", text, skipNul = FALSE)
```

Let's use `read.table`. Here we'll read the Titanic dataset from the University of Colorado Denver. Eventhough it off the Internet, the same use applies to your hard-drive. We set the separator parameter of `\t` for tabular and request to consider the first row as headers:

```
Titanic_dataset <- read.table('http://math.ucdenver.edu/RTutorial/titanic.txt', se
p='\t', header=TRUE)
head(Titanic_dataset)
```

```
##                               Name PClass   Age    Sex
## 1             Allen, Miss Elisabeth Walton    1st 29.00 female
## 2             Allison, Miss Helen Loraine    1st  2.00 female
## 3      Allison, Mr Hudson Joshua Creighton    1st 30.00   male
## 4 Allison, Mrs Hudson JC (Bessie Waldo Daniels)    1st 25.00 female
## 5             Allison, Master Hudson Trevor    1st  0.92   male
## 6             Anderson, Mr Harry    1st 47.00   male
## Survived
## 1      1
## 2      0
## 3      0
## 4      0
## 5      1
## 6      1
```

Let's see if the function figured out the data types contained in the Titanic dataset:

```
str(Titanic_dataset)
```

```
## 'data.frame':    1313 obs. of  5 variables:
## $ Name      : Factor w/ 1310 levels "Abbing, Mr Anthony",...: 22 25 26 27 24 31 4
5 46 50 54 ...
## $ PClass    : Factor w/ 3 levels "1st","2nd","3rd": 1 1 1 1 1 1 1 1 1 1 ...
## $ Age       : num  29 2 30 25 0.92 47 63 39 58 71 ...
## $ Sex       : Factor w/ 2 levels "female","male": 1 1 2 1 2 2 1 2 1 2 ...
## $ Survived: int   1 0 0 0 1 1 1 0 1 0 ...
```

Not bad, but feature `Names` shouldn't be a factor but a character. In this case, we could change it after the case or force the reader to make all text a `character` instead of `factor` by setting the `stringsAsFactors` to `FALSE`:

```
Titanic_dataset <- read.table('http://math.ucdenver.edu/RTutorial/titanic.txt', se
p='\t', header=TRUE, stringsAsFactors=FALSE)
str(Titanic_dataset)
```

```
## 'data.frame':    1313 obs. of  5 variables:
## $ Name      : chr  "Allen, Miss Elisabeth Walton" "Allison, Miss Helen Loraine"
"Allison, Mr Hudson Joshua Creighton" "Allison, Mrs Hudson JC (Bessie Waldo Daniel
s)" ...
## $ PClass    : chr  "1st" "1st" "1st" "1st" ...
## $ Age       : num  29 2 30 25 0.92 47 63 39 58 71 ...
## $ Sex       : chr  "female" "female" "male" "female" ...
## $ Survived: int   1 0 0 0 1 1 1 0 1 0 ...
```

You can also pass custom column names directly to the `read.table` function:

```
actg320_colnames <- c('id','time','censor','time_d','censor_d','treatment','treatm
ent_group','strat2','sex','raceth','ivdrug','hemophil','karnof','cd4','priorzd
v','age')
actg320 <- read.table('https://www.umass.edu/statdata/statdata/data/actg320.dat',
col.names = actg320_colnames)
head(actg320)
```

```
##   id time censor time_d censor_d treatment treatment_group strat2 sex
## 1  1  189      0   189      0          0          1      1  1
## 2  2  287      0   287      0          0          1      1  2
## 3  3  242      0   242      0          1          2      0  1
## 4  4  199      0   199      0          0          1      1  1
## 5  5  286      0   286      0          1          2      0  1
## 6  6  285      0   285      0          1          2      0  1
##   raceth ivdrug hemophil karnof   cd4 priorzdv age
## 1     1     1      0    100 169.0      39  34
## 2     2     1      0     90 149.5      15  34
## 3     1     1      1    100  23.5       9  20
## 4     1     1      0     90  46.0      53  48
## 5     1     3      0     90  10.0      12  46
## 6     1     1      0     70   0.0      24  51
```

```
dim(actg320)
```

```
## [1] 1151  16
```

Here's a quick look at `read.csv`. For more differences between these readers see: `read.table` (<http://www.inside-r.org/r-doc/utils/read.table>)

```
Iris_dataset <- read.csv('http://archive.ics.uci.edu/ml/machine-learning-database
s/iris/bezdekIris.data', header=FALSE)
head(Iris_dataset)
```

```
##   V1  V2  V3  V4      V5
## 1 5.1 3.5 1.4 0.2 Iris-setosa
## 2 4.9 3.0 1.4 0.2 Iris-setosa
## 3 4.7 3.2 1.3 0.2 Iris-setosa
## 4 4.6 3.1 1.5 0.2 Iris-setosa
## 5 5.0 3.6 1.4 0.2 Iris-setosa
## 6 5.4 3.9 1.7 0.4 Iris-setosa
```

Before we move to more sophisticated readers, let's build a simple data frame to work with. We'll create a very small data set made of the following data types: integers, factors, doubles, and dates.

```
mix_dataset <- data.frame(
  id=c(10,20,30,40,50),
  gender=c('male','female','female','male','female'),
  some_date=c('2012-01-12','2012-01-12','2012-12-01','2012-05-30','2
013-12-12'),
  value=c(12.34, 32.2, 24.3, 83.1, 8.32),
  outcome=c(1,1,0,0,0))

write.csv(mix_dataset, 'mix_dataset.csv', row.names = FALSE)
```

We'll save it to your current working directory and read it back again. Let's look at it using `read.csv`:

```
mix_dataset <- read.csv('mix_dataset.csv', stringsAsFactors = FALSE)
str(mix_dataset)
```

```
## 'data.frame':    5 obs. of  5 variables:
## $ id          : int  10 20 30 40 50
## $ gender      : chr  "male" "female" "female" "male" ...
## $ some_date   : chr  "2012-01-12" "2012-01-12" "2012-12-01" "2012-05-30" ...
## $ value       : num  12.34 32.2 24.3 83.1 8.32
## $ outcome     : int   1 1 0 0 0
```

The point to note here is the our date field `some_date` was read as a character string using `read.csv`.

## Heavy-duty Readers

### `readr`

Let's look at some readers that aren't part of the base package. `readr` {`readr`} (<https://cran.r-project.org/web/packages/readr/README.html>) is a relatively new package maintained by Hadley Wickham. It does a great job at inferring data types and is fast!

As you can see, `some_date` is correctly cast as `Date`, and it does pick up two integer fields:

```
# install.packages('readr')
library(readr)

mix_dataset <- read_csv('mix_dataset.csv')
str(mix_dataset)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    5 obs. of  5 variables:
## $ id      : int  10 20 30 40 50
## $ gender   : chr  "male" "female" "female" "male" ...
## $ some_date: Date, format: "2012-01-12" "2012-01-12" ...
## $ value    : num  12.34 32.2 24.3 83.1 8.32
## $ outcome  : int   1 1 0 0 0
```

If you know in advance the data types in a data set, you can pass it along to the function to save it time and processing. In the first read, we force `some_date` to character, in the read we force the `id` field to be a numeric instead of an integer.

```
mix_dataset <- read_csv('mix_dataset.csv', col_types='nccnn')
str(mix_dataset)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    5 obs. of  5 variables:
## $ id      : num  10 20 30 40 50
## $ gender   : chr  "male" "female" "female" "male" ...
## $ some_date: chr  "2012-01-12" "2012-01-12" "2012-12-01" "2012-05-30" ...
## $ value    : num  12.34 32.2 24.3 83.1 8.32
## $ outcome  : num   1 1 0 0 0
```

```
mix_dataset <- read_csv('mix_dataset.csv', col_types='ncDni')
str(mix_dataset)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    5 obs. of  5 variables:
## $ id      : num  10 20 30 40 50
## $ gender   : chr  "male" "female" "female" "male" ...
## $ some_date: Date, format: "2012-01-12" "2012-01-12" ...
## $ value    : num  12.34 32.2 24.3 83.1 8.32
## $ outcome  : int   1 1 0 0 0
```

## fread

`fread {data.table}` (<http://www.inside-r.org/packages/cran/data.table/docs/fread>) is the fastest of the bunch presented so far. Set parameter `data.table = FALSE` to return a data frame:

```
# install.packages('data.table')
library(data.table)
mix_dataset <- fread('mix_dataset.csv', showProgress=TRUE, data.table=FALSE)
str(mix_dataset)
```

```
## 'data.frame':    5 obs. of  5 variables:
## $ id      : int  10 20 30 40 50
## $ gender   : chr  "male" "female" "female" "male" ...
## $ some_date: chr  "2012-01-12" "2012-01-12" "2012-12-01" "2012-05-30" ...
## $ value    : num  12.34 32.2 24.3 83.1 8.32
## $ outcome  : int   1 1 0 0 0
```

If you only want a subset of columns, you can choose them directly in the `fread` command using the `select` parameter (similarly you can use the `drop` parameter to remove features):

```
mix_dataset <- fread('mix_dataset.csv', data.table=FALSE, select = c('value', 'outcome'))
head(mix_dataset)
```

```
##   value outcome
## 1 12.34        1
## 2 32.20        1
## 3 24.30        0
## 4 83.10        0
## 5  8.32        0
```

For reference, if you want to load data from an Excel spreadsheet, here are some popular libraries: `XLConnect` (<https://cran.r-project.org/web/packages/XLConnect/index.html>), `openxlsx` (<https://cran.r-project.org/web/packages/openxlsx/index.html>), `readxl` (<https://cran.r-project.org/web/packages/readxl/index.html>).

## Pipeline Check

So far we don't really need to create any custom code, simply use `read_csv` or `fread` if you have large data and/or complex data types, otherwise stick to `read.csv`. If this is new data, use the `readLines` function:

```
path_and_file_name <- 'https://cran.r-project.org/src/base/README'
print(readLines(path_and_file_name, n=5))
```

```
## [1] ""
## [2] "\t\t\tTHE BASIC R README"
## [3] ""
## [4] ""
## [5] "\t (See \"doc/FAQ\" and \"doc/RESOURCES\" for more detailed information"
```